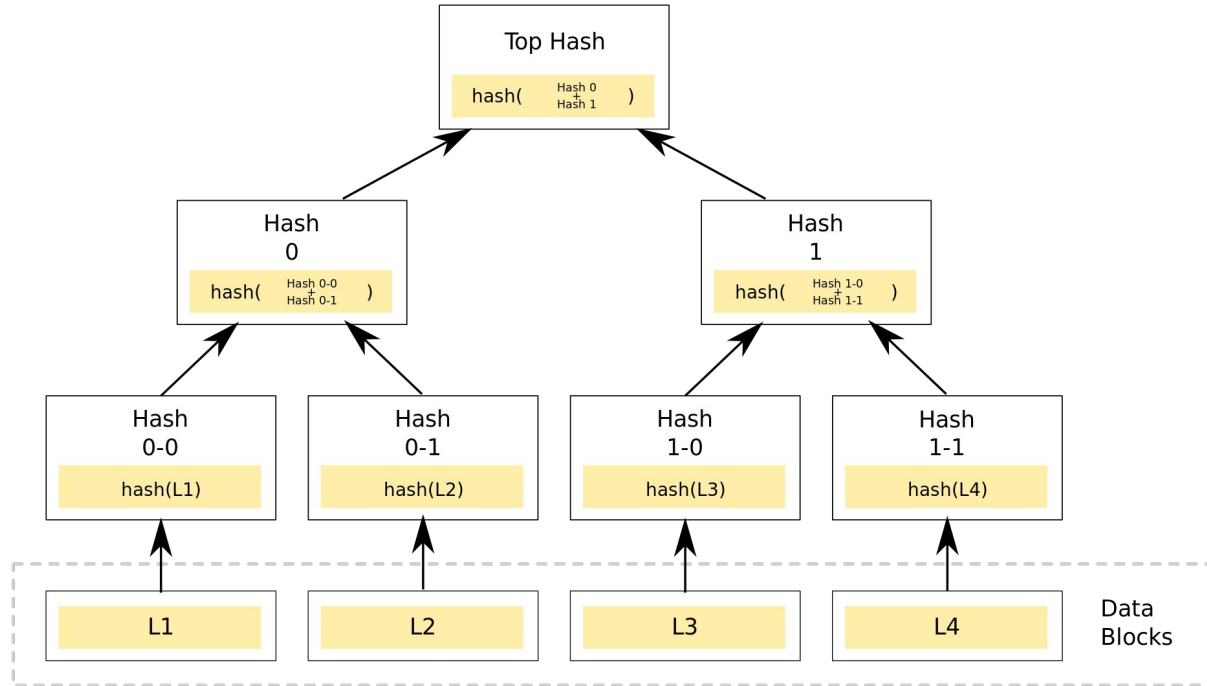


Concurrent Merkle Trees

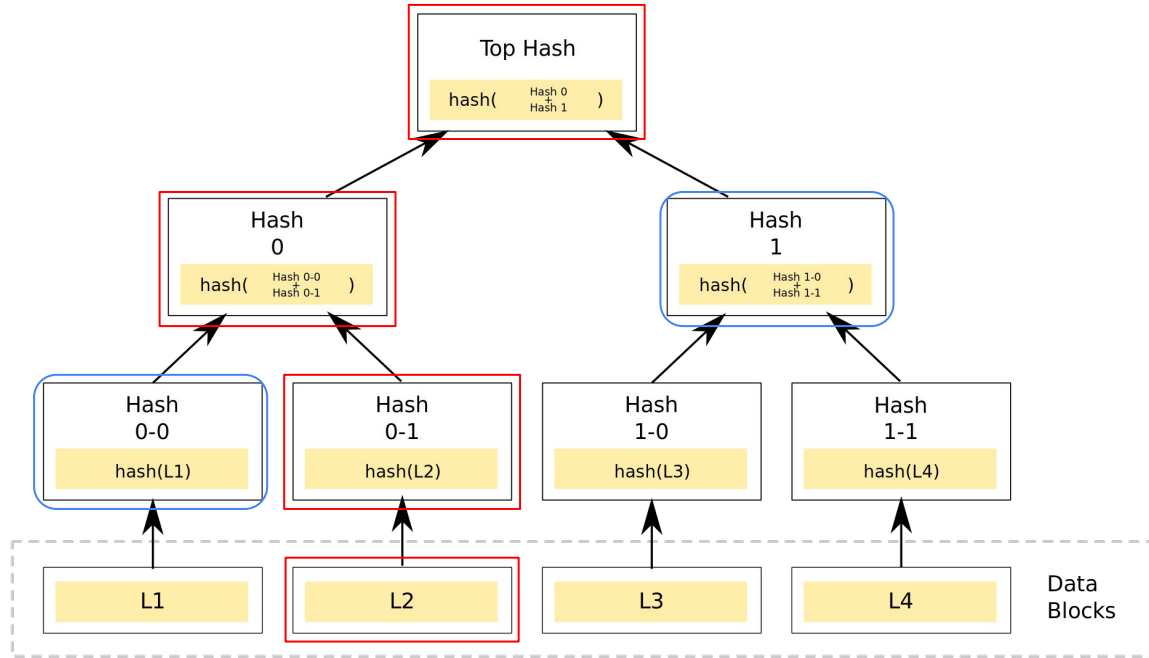
Gautam Singh (CS21BTECH11018)
K Vivek Kumar (CS21BTECH11026)

Merkle Trees



Merkle Trees: Updating a Leaf

- To update node L2, we have to update all nodes in red (Hash 0-1, Hash 0, Top Hash)
- Merkle proof is in blue (Hash 0-0, Hash 1).
 - Sometimes called “co-path”.
 - Used to verify the existence of a leaf by rebuilding the merkle root and checking equality.
- Both update and verify are of logarithmic complexity.



Concurrency conditions required for a continuously updating Merkle Tree

1. Concurrency in appending a data leaf

Potential Issues: Order of receiving data leaves

2. Concurrency in correcting a data leaf and involved hash nodes

Potential Issues: Delay in between identifying the false data leaf and updating the involved hash nodes

Concurrent Merkle Trees: Naive Implementation (C/C++)

```
for (unsigned int l = 0, n = 0; l < blocks.size(); l = l + 2, n++) {
    if (l != blocks.size() - 1) { // checks for adjacent block
        nodes[n] = new Node(md5(blocks[l]->hash_val + blocks[l+1]->hash_val));
        nodes[n]->mom = blocks[l]; // assign children
        nodes[n]->dad = blocks[l+1];
    } else {
        nodes[n] = blocks[l];
    }
}

//std::cout << "\n";
blocks = nodes;
nodes.clear();
}

this->root = blocks[0];
}
```

Fig. 3: The serial implementation used in defining the Merkle tree.

```
#pragma omp parallel
{
    #pragma omp for ordered
    for (unsigned int l = 0, n = 0; l < blocks.size(); l = l + 2, n++) {
        if (l != blocks.size() - 1) { // checks for adjacent block
            nodes[n] = new Node(md5(blocks[l]->hash_val + blocks[l+1]->hash_val));
            nodes[n]->mom = blocks[l]; // assign children
            nodes[n]->dad = blocks[l+1];
        } else {
            nodes[n] = blocks[l];
        }
    }

    //std::cout << "\n";
    blocks = nodes;
    nodes.clear();
}

this->root = blocks[0];
}
```

Fig. 4: The parallel implementation used in defining the Merkle tree.

Concurrent Merkle Trees: Using Changelogs (Rust/TS)

- How to validate the proofs?
 - Use changelogs!
- After previous updates, keep a record of the old values in the Merkle tree.
- For new update, search for the (old) proof in the changelogs and build from there.




Figure 3: The **cyan** path and the **orange** path modify the same tree, but once one of the changes is locked in, the proof to the other change will be invalid.

Concurrent Merkle Trees: Using Changelogs (Rust/TS)

Hyperparameter of changelog
buffer size (K).

- Time complexity depends on additional factor.
- Storage overheads.
- Still not used in Solana ecosystem.







Figure 3: The **cyan** path and the **orange** path modify the same tree, but once one of the changes is locked in, the proof to the other change will be invalid.

Concurrent Merkle Trees: Fine-Grained Locking (Go/Python)



Tree split among worker nodes (one subtree per AWS EC2 instance), race conditions are ONLY at conflict nodes.





Concurrent Merkle Trees: Fine-Grained Locking (Go/Python)

- Each update edge traversed only once.
 - Threads that arrive early at a conflict node return and are reused elsewhere.
- Lazy locking used here.
 - Threads lock only when they reach a conflict node.
- Overhead of maintaining conflict set.
 - We attempt to address this later!

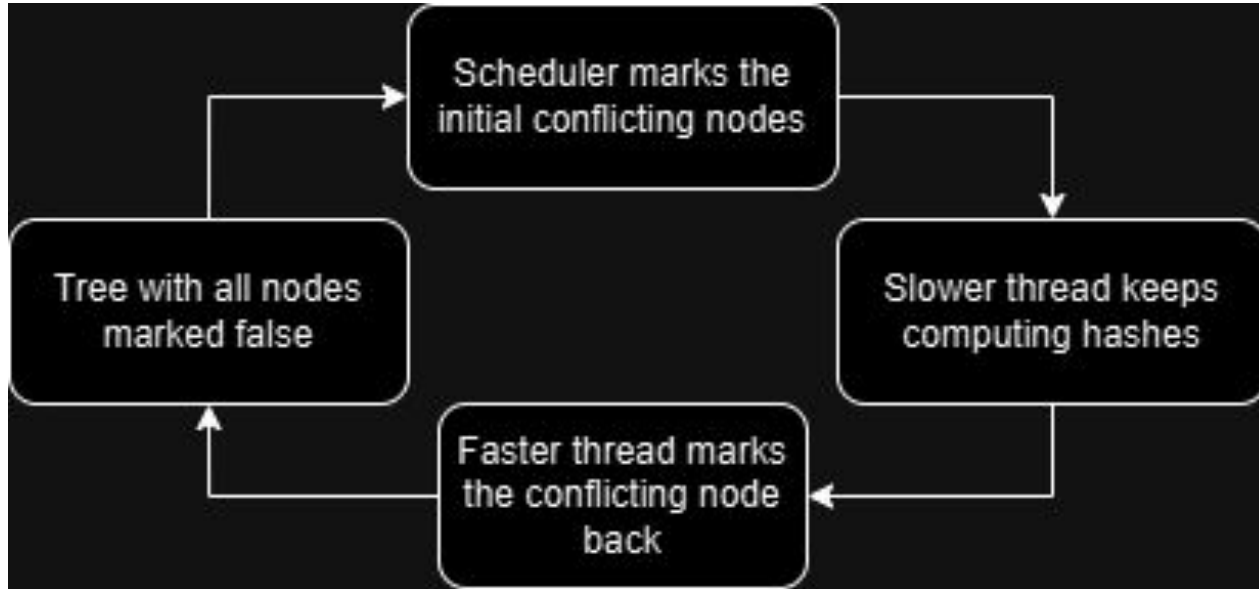
```
procedure updateTree(Transaction T)  
  while parent_exists() do  
    if parent in Tree.conflicts then  
      Tree.conflicts(parent).acquire_lock()  
      defer Tree.conflicts(parent).release_lock()  
      if parent has not been visited then  
        Mark Node as visited  
      End Thread  
    else  
      Process Update  
    end if  
  else  
    Process Update  
  end if  
end while  
end procedure
```

Our Proposition


- **Setting: Scheduler decides operations in “read” and “write” phases.**
 - Similar to Angela
 - Reads: Verifications
 - Writes: Updates
- **Use barriers at conflict nodes.**
 - Unlike traditional barriers, threads don’t wait but instead simply return.
 - Avoid lock acquire/release overheads as in Angela.
- **For binary merkle trees, the barrier is effectively CAS.**
 - Scheduler marks conflict nodes.
 - Threads atomically unmark them.



Balancing the nodes markings



Thread life visualization



Algorithm

Algorithm 2 Algorithm for Single Node Update

```
1: procedure UPDATE(Update  $u$ , Tree  $T$ )
2:    $node \leftarrow u.leaf$ 
3:   while  $node \neq null$  do
4:      $node \leftarrow \text{COMPUTEHASH}(node)$ 
5:     if  $node.parent = null$  then
6:       return ▷ Root hash computed
7:     end if
8:      $node \leftarrow node.parent$ 
9:      $marked \leftarrow \text{CASMARK}(node, true, false)$ 
10:    if  $marked$  then
11:      return ▷ Conflict detected
12:    end if
13:  end while
14: end procedure
```

Observations/Learnings

- Concurrency in Merkle Trees is not easy
 - Inherent dependencies
 - Need to process updates from bottom up
- Deficiencies of read-write scheduling
 - Overhead on scheduler
 - No concurrency between update and verify operations.
- What next?
 - Benchmark our proposition (implement one subtree per thread instead of one operation per thread).
 - Can we have a concurrent update and verify without significant overheads for lightweight blockchain clients (such changelog buffers in case of Solana)?
 - Better work splitting? One thread per subtree?